

# Multi-class\_Classification

October 13, 2025

## 1 Multi-class Classification

Estimated time needed: **30** mins

In this lab, you will learn the different strategies of Multi-class classification and implement the same on a real-world dataset.

### 1.1 Objectives

After completing this lab you will be able to:

1. Understand the use of one-hot encoding for categorical variables.
2. Implement logistic regression for multi-class classification using **One-vs-All (OvA)** and **One-vs-One (OvO)** strategies.
3. Evaluate model performance using appropriate metrics.

### 1.2 Import Necessary Libraries

First, to ensure the availability of the required libraries, execute the cell below.

```
[1]: !pip install numpy==2.2.0
!pip install pandas==2.2.3
!pip install scikit-learn==1.6.0
!pip install matplotlib==3.9.3
!pip install seaborn==0.13.2
```

```
Collecting numpy==2.2.0
  Downloading
numpy-2.2.0-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata
(62 kB)
  Downloading
numpy-2.2.0-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (16.1 MB)
    16.1/16.1 MB
 170.8 MB/s eta 0:00:00
Installing collected packages: numpy
Successfully installed numpy-2.2.0
Collecting pandas==2.2.3
  Downloading
pandas-2.2.3-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata
(89 kB)
```

```
Requirement already satisfied: numpy>=1.26.0 in /opt/conda/lib/python3.12/site-packages (from pandas==2.2.3) (2.2.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /opt/conda/lib/python3.12/site-packages (from pandas==2.2.3) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /opt/conda/lib/python3.12/site-packages (from pandas==2.2.3) (2024.2)
Collecting tzdata>=2022.7 (from pandas==2.2.3)
  Downloading tzdata-2025.2-py2.py3-none-any.whl.metadata (1.4 kB)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.12/site-packages (from python-dateutil>=2.8.2->pandas==2.2.3) (1.17.0)
Downloading
pandas-2.2.3-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (12.7 MB)
  12.7/12.7 MB
134.1 MB/s eta 0:00:00
Downloading tzdata-2025.2-py2.py3-none-any.whl (347 kB)
Installing collected packages: tzdata, pandas
Successfully installed pandas-2.2.3 tzdata-2025.2
Collecting scikit-learn==1.6.0
  Downloading scikit_learn-1.6.0-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (18 kB)
Requirement already satisfied: numpy>=1.19.5 in /opt/conda/lib/python3.12/site-packages (from scikit-learn==1.6.0) (2.2.0)
Collecting scipy>=1.6.0 (from scikit-learn==1.6.0)
  Downloading scipy-1.16.2-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.whl.metadata (62 kB)
Collecting joblib>=1.2.0 (from scikit-learn==1.6.0)
  Downloading joblib-1.5.2-py3-none-any.whl.metadata (5.6 kB)
Collecting threadpoolctl>=3.1.0 (from scikit-learn==1.6.0)
  Downloading threadpoolctl-3.6.0-py3-none-any.whl.metadata (13 kB)
Downloading
scikit_learn-1.6.0-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (13.1 MB)
  13.1/13.1 MB
144.5 MB/s eta 0:00:00
Downloading joblib-1.5.2-py3-none-any.whl (308 kB)
Downloading
scipy-1.16.2-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (35.7 MB)
  35.7/35.7 MB
231.0 MB/s eta 0:00:00
Downloading threadpoolctl-3.6.0-py3-none-any.whl (18 kB)
Installing collected packages: threadpoolctl, scipy, joblib, scikit-learn
Successfully installed joblib-1.5.2 scikit-learn-1.6.0 scipy-1.16.2
threadpoolctl-3.6.0
Collecting matplotlib==3.9.3
  Downloading matplotlib-3.9.3-cp312-cp312-
```

```

manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (11 kB)
Collecting contourpy>=1.0.1 (from matplotlib==3.9.3)
    Downloading contourpy-1.3.3-cp312-cp312-
manylinux_2_27_x86_64.manylinux_2_28_x86_64.whl.metadata (5.5 kB)
Collecting cycler>=0.10 (from matplotlib==3.9.3)
    Downloading cycler-0.12.1-py3-none-any.whl.metadata (3.8 kB)
Collecting fonttools>=4.22.0 (from matplotlib==3.9.3)
    Downloading fonttools-4.60.0-cp312-cp312-
manylinux1_x86_64.manylinux2014_x86_64.manylinux_2_17_x86_64.manylinux_2_5_x86_6
4.whl.metadata (111 kB)
Collecting kiwisolver>=1.3.1 (from matplotlib==3.9.3)
    Downloading kiwisolver-1.4.9-cp312-cp312-
manylinux2014_x86_64.manylinux_2_17_x86_64.whl.metadata (6.3 kB)
Requirement already satisfied: numpy>=1.23 in /opt/conda/lib/python3.12/site-
packages (from matplotlib==3.9.3) (2.2.0)
Requirement already satisfied: packaging>=20.0 in
/opt/conda/lib/python3.12/site-packages (from matplotlib==3.9.3) (24.2)
Collecting pillow>=8 (from matplotlib==3.9.3)
    Downloading pillow-11.3.0-cp312-cp312-
manylinux_2_27_x86_64.manylinux_2_28_x86_64.whl.metadata (9.0 kB)
Collecting pyparsing>=2.3.1 (from matplotlib==3.9.3)
    Downloading pyparsing-3.2.5-py3-none-any.whl.metadata (5.0 kB)
Requirement already satisfied: python-dateutil>=2.7 in
/opt/conda/lib/python3.12/site-packages (from matplotlib==3.9.3) (2.9.0.post0)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.12/site-
packages (from python-dateutil>=2.7->matplotlib==3.9.3) (1.17.0)
Downloading
matplotlib-3.9.3-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (8.3
MB)
    8.3/8.3 MB
142.1 MB/s eta 0:00:00
Downloading
contourpy-1.3.3-cp312-cp312-manylinux_2_27_x86_64.manylinux_2_28_x86_64.whl (362
kB)
Downloading cycler-0.12.1-py3-none-any.whl (8.3 kB)
Downloading fonttools-4.60.0-cp312-cp312-
manylinux1_x86_64.manylinux2014_x86_64.manylinux_2_17_x86_64.manylinux_2_5_x86_6
4.whl (4.9 MB)
    4.9/4.9 MB
143.7 MB/s eta 0:00:00
Downloading
kiwisolver-1.4.9-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (1.5
MB)
    1.5/1.5 MB
89.4 MB/s eta 0:00:00
Downloading
pillow-11.3.0-cp312-cp312-manylinux_2_27_x86_64.manylinux_2_28_x86_64.whl (6.6
MB)

```

6.6/6.6 MB

```
100.7 MB/s eta 0:00:00
Downloading pyparsing-3.2.5-py3-none-any.whl (113 kB)
Installing collected packages: pyparsing, pillow, kiwisolver, fonttools, cycler,
contourpy, matplotlib
Successfully installed contourpy-1.3.3 cycler-0.12.1 fonttools-4.60.0
kiwisolver-1.4.9 matplotlib-3.9.3 pillow-11.3.0 pyparsing-3.2.5
Collecting seaborn==0.13.2
    Downloading seaborn-0.13.2-py3-none-any.whl.metadata (5.4 kB)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in
/opt/conda/lib/python3.12/site-packages (from seaborn==0.13.2) (2.2.0)
Requirement already satisfied: pandas>=1.2 in /opt/conda/lib/python3.12/site-
packages (from seaborn==0.13.2) (2.2.3)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in
/opt/conda/lib/python3.12/site-packages (from seaborn==0.13.2) (3.9.3)
Requirement already satisfied: contourpy>=1.0.1 in
/opt/conda/lib/python3.12/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn==0.13.2) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /opt/conda/lib/python3.12/site-
packages (from matplotlib!=3.6.1,>=3.4->seaborn==0.13.2) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/opt/conda/lib/python3.12/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn==0.13.2) (4.60.0)
Requirement already satisfied: kiwisolver>=1.3.1 in
/opt/conda/lib/python3.12/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn==0.13.2) (1.4.9)
Requirement already satisfied: packaging>=20.0 in
/opt/conda/lib/python3.12/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn==0.13.2) (24.2)
Requirement already satisfied: pillow>=8 in /opt/conda/lib/python3.12/site-
packages (from matplotlib!=3.6.1,>=3.4->seaborn==0.13.2) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/opt/conda/lib/python3.12/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn==0.13.2) (3.2.5)
Requirement already satisfied: python-dateutil>=2.7 in
/opt/conda/lib/python3.12/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn==0.13.2) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /opt/conda/lib/python3.12/site-
packages (from pandas>=1.2->seaborn==0.13.2) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /opt/conda/lib/python3.12/site-
packages (from pandas>=1.2->seaborn==0.13.2) (2025.2)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.12/site-
packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn==0.13.2)
(1.17.0)
Downloading seaborn-0.13.2-py3-none-any.whl (294 kB)
Installing collected packages: seaborn
Successfully installed seaborn-0.13.2
```

Now, import the necessary libraries for data processing, model training, and evaluation.

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.multiclass import OneVsOneClassifier
from sklearn.metrics import accuracy_score

import warnings
warnings.filterwarnings('ignore')
```

### 1.3 About the dataset

The data set being used for this lab is the “Obesity Risk Prediction” data set publically available on UCI Library under the CCA 4.0 license. The data set has 17 attributes in total along with 2,111 samples.

The attributes of the dataset are described below.

Variable Name

Type

Description

Gender

Categorical

Age

Continuous

Height

Continuous

Weight

Continuous

family\_history\_with\_overweight

Binary

Has a family member suffered or suffers from overweight?

FAVC

Binary

Do you eat high caloric food frequently?

FCVC

Integer

Do you usually eat vegetables in your meals?

NCP

Continuous

How many main meals do you have daily?

CAEC

Categorical

Do you eat any food between meals?

SMOKE

Binary

Do you smoke?

CH2O

Continuous

How much water do you drink daily?

SCC

Binary

Do you monitor the calories you eat daily?

FAF

Continuous

How often do you have physical activity?

TUE

Integer

How much time do you use technological devices such as cell phone, videogames, television, computer and others?

CALC

Categorical

How often do you drink alcohol?

MTRANS

Categorical

Which transportation do you usually use?

NObeyesdad

Categorical

Obesity level

### 1.3.1 Load the dataset

Load the data set by executing the code cell below.

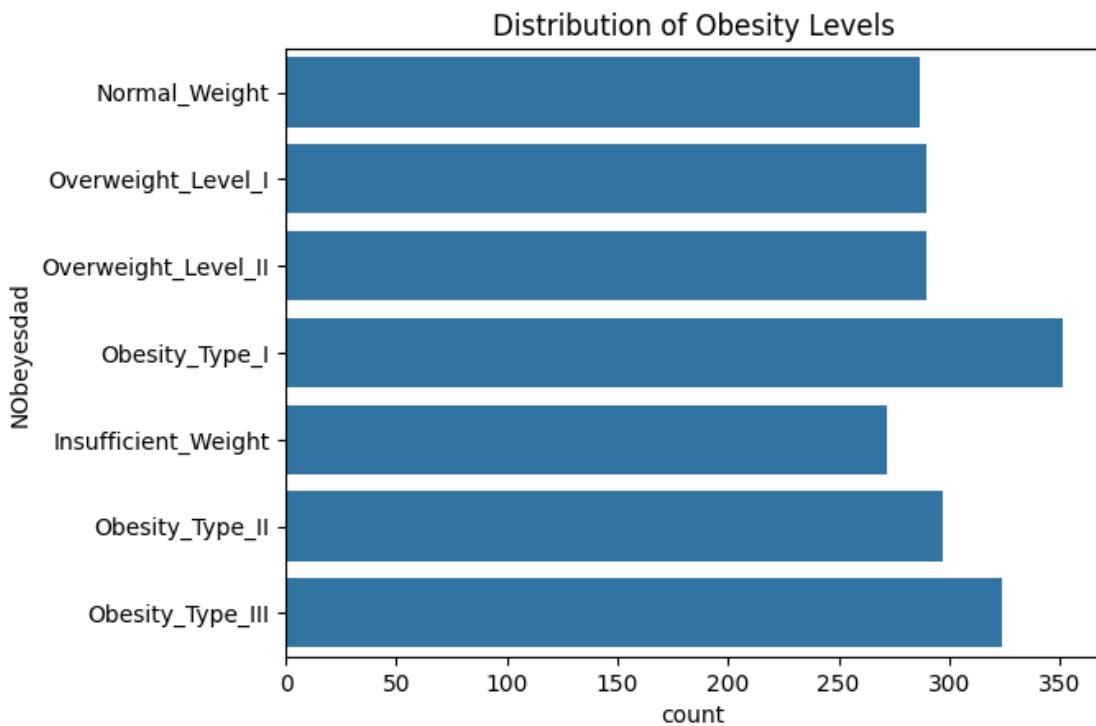
```
[3]: file_path = "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/  
↳GkDzb7bWrtyGXdP0fk6CIg/Obesity-level-prediction-dataset.csv"  
data = pd.read_csv(file_path)  
data.head()
```

```
[3]:   Gender    Age  Height  Weight family_history_with_overweight FAVC  FCVC  \  
0  Female  21.0    1.62    64.0                               yes  no  2.0  
1  Female  21.0    1.52    56.0                               yes  no  3.0  
2   Male  23.0    1.80    77.0                               yes  no  2.0  
3   Male  27.0    1.80    87.0                              no  no  3.0  
4   Male  22.0    1.78    89.8                              no  no  2.0  
  
      NCP        CAEC SMOKE  CH20  SCC  FAF  TUE          CALC  \  
0  3.0  Sometimes  no  2.0  no  0.0  1.0          no  
1  3.0  Sometimes  yes  3.0  yes  3.0  0.0  Sometimes  
2  3.0  Sometimes  no  2.0  no  2.0  1.0  Frequently  
3  3.0  Sometimes  no  2.0  no  2.0  0.0  Frequently  
4  1.0  Sometimes  no  2.0  no  0.0  0.0  Sometimes  
  
      MTRANS          NObeyesdad  
0  Public_Transportation  Normal_Weight  
1  Public_Transportation  Normal_Weight  
2  Public_Transportation  Normal_Weight  
3           Walking  Overweight_Level_I  
4  Public_Transportation  Overweight_Level_II
```

## 1.4 Exploratory Data Analysis

Visualize the distribution of the target variable to understand the class balance.

```
[4]: # Distribution of target variable  
sns.countplot(y='NObeyesdad', data=data)  
plt.title('Distribution of Obesity Levels')  
plt.show()
```



This shows that the dataset is fairly balanced and does not require any special attention in terms of biased training.

#### 1.4.1 Exercise 1

Check for null values, and display a summary of the dataset (use `.info()` and `.describe()` methods).

```
[10]: # your code here
#check Null values
print(data.isnull().sum())
print(data.info())
print(data.describe())
```

Gender	0
Age	0
Height	0
Weight	0
family_history_with_overweight	0
FAVC	0
FCVC	0
NCP	0
CAEC	0
SMOKE	0
CH2O	0

```

SCC          0
FAF          0
TUE          0
CALC         0
MTRANS        0
NObeyesdad   0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2111 entries, 0 to 2110
Data columns (total 17 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Gender             2111 non-null    object  
 1   Age                2111 non-null    float64 
 2   Height             2111 non-null    float64 
 3   Weight              2111 non-null    float64 
 4   family_history_with_overweight  2111 non-null    object  
 5   FAVC               2111 non-null    object  
 6   FCVC               2111 non-null    float64 
 7   NCP                2111 non-null    float64 
 8   CAEC               2111 non-null    object  
 9   SMOKE              2111 non-null    object  
 10  CH20               2111 non-null    float64 
 11  SCC                2111 non-null    object  
 12  FAF                2111 non-null    float64 
 13  TUE                2111 non-null    float64 
 14  CALC               2111 non-null    object  
 15  MTRANS              2111 non-null    object  
 16  NObeyesdad         2111 non-null    object  
dtypes: float64(8), object(9)
memory usage: 280.5+ KB
None
      Age      Height     Weight      FCVC      NCP \ 
count  2111.000000  2111.000000  2111.000000  2111.000000  2111.000000
mean    24.312600    1.701677    86.586058    2.419043    2.685628
std     6.345968    0.093305   26.191172    0.533927    0.778039
min    14.000000    1.450000    39.000000    1.000000    1.000000
25%   19.947192    1.630000    65.473343    2.000000    2.658738
50%   22.777890    1.700499    83.000000    2.385502    3.000000
75%   26.000000    1.768464   107.430682    3.000000    3.000000
max    61.000000    1.980000   173.000000    3.000000    4.000000

      CH20      FAF      TUE
count  2111.000000  2111.000000  2111.000000
mean    2.008011    1.010298    0.657866
std     0.612953    0.850592    0.608927
min    1.000000    0.000000    0.000000
25%   1.584812    0.124505    0.000000

```

```

50%      2.000000    1.000000    0.625350
75%      2.477420    1.666678    1.000000
max      3.000000    3.000000    2.000000

```

[Click here for the solution](#)

```

# Checking for null values
print(data.isnull().sum())

# Dataset summary
print(data.info())
print(data.describe())

```

Expected Output:

- Counts of null values for each column (likely zero for this dataset).
- Dataset info including column names, data types, and memory usage.
- Descriptive statistics for numerical columns.

## 1.5 Preprocessing the data

### 1.5.1 Feature scaling

Scale the numerical features to standardize their ranges for better model performance.

```
[14]: # Standardizing continuous numerical features
continuous_columns = data.select_dtypes(include=['float64']).columns.tolist()

scaler = StandardScaler()
scaled_features = scaler.fit_transform(data[continuous_columns])

# Converting to a DataFrame
scaled_df = pd.DataFrame(scaled_features, columns=scaler.
                           ↪get_feature_names_out(continuous_columns))

# Combining with the original dataset
scaled_data = pd.concat([data.drop(columns=continuous_columns), scaled_df], ↪
                           ↪axis=1)
data.sample(10)
```

```
[14]:   Gender      Age      Height      Weight family_history_with_overweight \
986  Female  28.583944  1.578560  65.522744                      yes
1357   Male  18.000000  1.784402 108.413119                      yes
647  Female  20.744839  1.667852  49.803921                      yes
782  Female  32.593129  1.721903  72.748903                      yes
166  Female  29.000000  1.740000  72.000000                      yes
303  Female  16.000000  1.570000  49.000000                     no
1703   Male  22.580038  1.849507 121.560938                      yes
157   Male  21.000000  1.670000  60.000000                      yes
232  Female  51.000000  1.590000  50.000000                      yes
```

2046	Female	25.895546	1.626179	110.074019	yes
------	--------	-----------	----------	------------	-----

	FAVC	FCVC	NCP	CAEC	SMOKE	CH20	SCC	FAF	\
986	no	2.000000	2.283673	Sometimes	no	1.970622	no	1.863258	
1357	yes	2.000000	2.475444	Sometimes	no	2.655795	no	1.000000	
647	yes	2.977018	3.193671	Frequently	no	2.482933	no	2.000000	
782	yes	2.000000	3.000000	Sometimes	no	1.000000	no	0.000000	
166	no	3.000000	3.000000	Always	no	2.000000	no	0.000000	
303	yes	2.000000	4.000000	Always	no	2.000000	no	0.000000	
1703	yes	3.000000	2.272801	Sometimes	no	1.560064	no	0.796770	
157	yes	2.000000	3.000000	Frequently	no	2.000000	no	2.000000	
232	no	3.000000	3.000000	Sometimes	yes	3.000000	yes	2.000000	
2046	yes	3.000000	3.000000	Sometimes	no	1.967707	no	0.014370	

	TUE	CALC	MTRANS	NObeyesdad
986	1.581703	Sometimes	Public_Transportation	Overweight_Level_II
1357	0.169879	no	Public_Transportation	Obesity_Type_I
647	1.000000	no	Public_Transportation	Insufficient_Weight
782	1.339232	Sometimes	Automobile	Overweight_Level_I
166	0.000000	Sometimes	Automobile	Normal_Weight
303	1.000000	Sometimes	Public_Transportation	Normal_Weight
1703	0.000000	Sometimes	Public_Transportation	Obesity_Type_II
157	1.000000	Sometimes	Public_Transportation	Normal_Weight
232	0.000000	no	Public_Transportation	Normal_Weight
2046	0.434073	Sometimes	Public_Transportation	Obesity_Type_III

Standardization of data is important to better define the decision boundaries between classes by making sure that the feature variations are in similar scales. The data is now ready to be used for training and testing.

### 1.5.2 One-hot encoding

Convert categorical variables into numerical format using one-hot encoding.

```
[15]: # Identifying categorical columns
categorical_columns = scaled_data.select_dtypes(include=['object']).columns.
    ↪tolist()
categorical_columns.remove('NObeyesdad') # Exclude target column

# Applying one-hot encoding
encoder = OneHotEncoder(sparse_output=False, drop='first')
encoded_features = encoder.fit_transform(scaled_data[categorical_columns])

# Converting to a DataFrame
encoded_df = pd.DataFrame(encoded_features, columns=encoder.
    ↪get_feature_names_out(categorical_columns))

# Combining with the original dataset
```

```
prepped_data = pd.concat([scaled_data.drop(columns=categorical_columns),
                           encoded_df], axis=1)
```

You will observe that all the categorical variables have now been modified to one-hot encoded features. This increases the overall number of fields to 24.

### 1.5.3 Encode the target variable

```
[16]: # Encoding the target variable
prepped_data['NObeyesdad'] = prepped_data['NObeyesdad'].astype('category').cat.
    codes
prepped_data.head()
```

```
[16]: NObeyesdad      Age     Height     Weight      FCVC      NCP      CH20 \
0           1 -0.522124 -0.875589 -0.862558 -0.785019  0.404153 -0.013073
1           1 -0.522124 -1.947599 -1.168077  1.088342  0.404153  1.618759
2           1 -0.206889  1.054029 -0.366090 -0.785019  0.404153 -0.013073
3           5  0.423582  1.054029  0.015808  1.088342  0.404153 -0.013073
4           6 -0.364507  0.839627  0.122740 -0.785019 -2.167023 -0.013073

      FAF      TUE  Gender_Male ... CAEC_no  SMOKE_yes  SCC_yes \
0 -1.188039  0.561997        0.0 ...        0.0        0.0        0.0
1  2.339750 -1.080625        0.0 ...        0.0        1.0        1.0
2  1.163820  0.561997        1.0 ...        0.0        0.0        0.0
3  1.163820 -1.080625        1.0 ...        0.0        0.0        0.0
4 -1.188039 -1.080625        1.0 ...        0.0        0.0        0.0

      CALC_Frequently  CALC_Sometimes  CALC_no  MTRANS_Bike  MTRANS_Motorbike \
0            0.0          0.0          1.0        0.0            0.0
1            0.0          1.0          0.0        0.0            0.0
2            1.0          0.0          0.0        0.0            0.0
3            1.0          0.0          0.0        0.0            0.0
4            0.0          1.0          0.0        0.0            0.0

      MTRANS_Public_Transportation  MTRANS_Walking
0                  1.0            0.0
1                  1.0            0.0
2                  1.0            0.0
3                  0.0            1.0
4                  1.0            0.0

[5 rows x 24 columns]
```

#### 1.5.4 Separate the input and target data

```
[17]: # Preparing final dataset  
X = prepped_data.drop('NObeyesdad', axis=1)  
y = prepped_data['NObeyesdad']
```

### 1.6 Model training and evaluation

#### 1.6.1 Splitting the data set

Split the data into training and testing subsets.

```
[18]: # Splitting data  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42, stratify=y)
```

#### 1.6.2 Logistic Regression with One-vs-All

In the One-vs-All approach:

- The algorithm trains a single binary classifier for each class.
- Each classifier learns to distinguish a single class from all the others combined.
- If there are k classes, k classifiers are trained.
- During prediction, the algorithm evaluates all classifiers on each input, and selects the class with the highest confidence score as the predicted class.

#### Advantages:

- Simpler and more efficient in terms of the number of classifiers (k)
- Easier to implement for algorithms that naturally provide confidence scores (e.g., logistic regression, SVM).

#### Disadvantages:

- Classifiers may struggle with class imbalance since each binary classifier must distinguish between one class and the rest.
- Requires the classifier to perform well even with highly imbalanced datasets, as the “all” group typically contains more samples than the “one” class.

Train a logistic regression model using the One-vs-All strategy and evaluate its performance.

```
[19]: # Training logistic regression model using One-vs-All (default)  
model_ova = LogisticRegression(multi_class='ovr', max_iter=1000)  
model_ova.fit(X_train, y_train)
```

```
[19]: LogisticRegression(max_iter=1000, multi_class='ovr')
```

You can now evaluate the accuracy of the trained model as a measure of its performance on unseen testing data.

```
[20]: # Predictions
y_pred_ova = model_ova.predict(X_test)

# Evaluation metrics for OvA
print("One-vs-All (OvA) Strategy")
print(f"Accuracy: {np.round(100*accuracy_score(y_test, y_pred_ova),2)}%")
```

One-vs-All (OvA) Strategy  
Accuracy: 76.12%

### 1.6.3 Logistic Regression with OvO

In the One-vs-One approach:

- \* The algorithm trains a binary classifier for every pair of classes in the dataset.
- \* If there are k classes, this results in  $k(k - 1)/2$  classifiers.
- \* Each classifier is trained to distinguish between two specific classes, ignoring the rest.
- \* During prediction, all classifiers are used, and a “voting” mechanism decides the final class by selecting the class that wins the majority of pairwise comparisons.

#### Advantages:

- Suitable for algorithms that are computationally expensive to train on many samples because each binary classifier deals with a smaller dataset (only samples from two classes).
- Can be more accurate in some cases since classifiers focus on distinguishing between two specific classes at a time.

#### Disadvantages:

- Computationally expensive for datasets with a large number of classes due to the large number of classifiers required.
- May lead to ambiguous predictions if voting results in a tie.

Train a logistic regression model using the One-vs-One (OvO) strategy and evaluate its performance.

```
[22]: # Training logistic regression model using One-vs-One
model_ovo = OneVsOneClassifier(LogisticRegression(max_iter=1000))
model_ovo.fit(X_train, y_train)
```

```
[22]: OneVsOneClassifier(estimator=LogisticRegression(max_iter=1000))
```

Evaluate the accuracy of the trained model as a measure of its performance on unseen testing data.

```
[24]: # Predictions
y_pred_ovo = model_ovo.predict(X_test)

# Evaluation metrics for OvO
print("One-vs-One (OvO) Strategy")
print(f"Accuracy: {np.round(100*accuracy_score(y_test, y_pred_ovo),2)}%")
```

One-vs-One (OvO) Strategy  
Accuracy: 92.2%

#### 1.6.4 Exercises

Q1. Experiment with different test sizes in the train\_test\_split method (e.g., 0.1, 0.3) and observe the impact on model performance.

```
[30]: # your code here
for test_size in [0.1,0.3]:
    X_train,X_test, y_train , y_test= train_test_split(X,y, test_size=test_size, random_state=42, stratify=y)
    model_ovo.fit(X_train,y_train)
    y_pred = model_ova.predict(X_test)
    print(f'Test size: {test_size}')
    print("Accuracy:", accuracy_score(y_test, y_pred))
```

Test size: 0.1  
Accuracy: 0.7641509433962265  
Test size: 0.3  
Accuracy: 0.7665615141955836

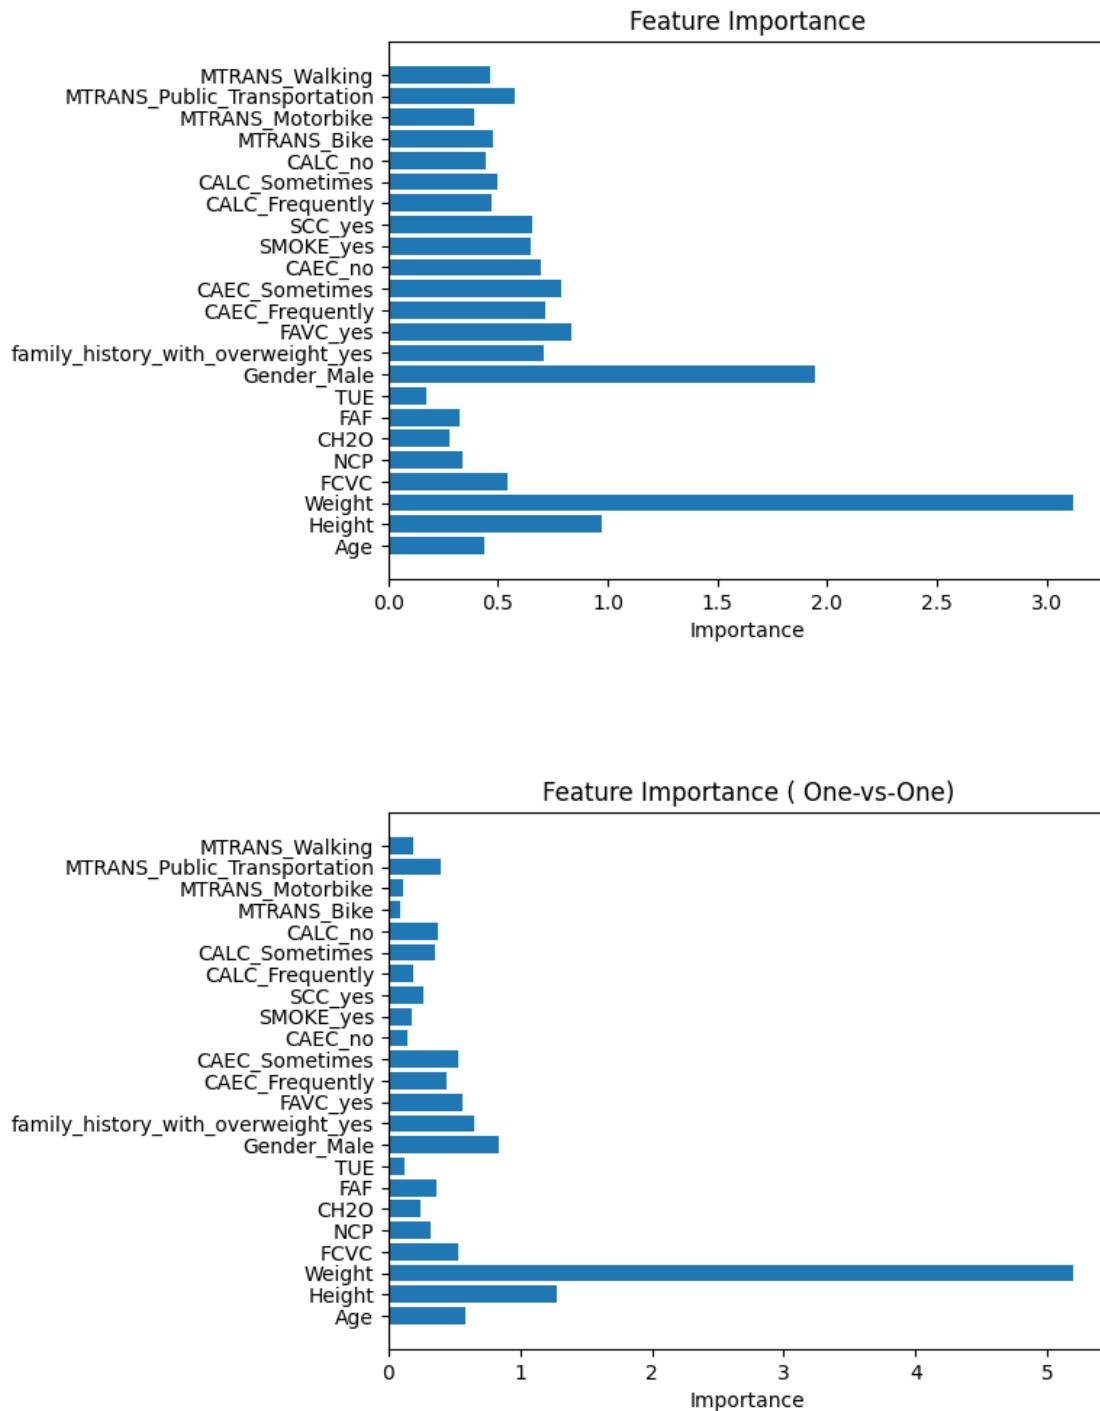
[Click here for the solution](#)

```
for test_size in [0.1, 0.3]:
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size, random_state=42)
    model_ovo.fit(X_train, y_train)
    y_pred = model_ovo.predict(X_test)
    print(f"Test Size: {test_size}")
    print("Accuracy:", accuracy_score(y_test, y_pred))
```

Q2. Plot a bar chart of feature importance using the coefficients from the One vs All logistic regression model. Also try for the One vs One model.

```
[35]: # your code here
#feature importance
feature_importance = np.mean(np.abs(model_ova.coef_), axis = 0)
plt.barh(X.columns, feature_importance)
plt.title("Feature Importance")
plt.xlabel("Importance")
plt.show()
# For One vs One model
# Collect all coefficients from each underlying binary classifier
coefs = np.array([est.coef_[0] for est in model_ovo.estimators_])

#Now take the Mean accross all the those classifiers
feature_importance = np.mean(np.abs(coefs), axis=0)
# plot feature importance
plt.barh(X.columns, feature_importance )
plt.title("Feature Importance ( One-vs-One)")
plt.xlabel("Importance")
plt.show()
```



Click here for the solution

```
# Feature importance
feature_importance = np.mean(np.abs(model_ova.coef_), axis=0)
plt.barh(X.columns, feature_importance)
```

```

plt.title("Feature Importance")
plt.xlabel("Importance")
plt.show()

# For One vs One model
# Collect all coefficients from each underlying binary classifier
coefs = np.array([est.coef_[0] for est in model_ovo.estimators_])

# Now take the mean across all those classifiers
feature_importance = np.mean(np.abs(coefs), axis=0)

# Plot feature importance
plt.barh(X.columns, feature_importance)
plt.title("Feature Importance (One-vs-One)")
plt.xlabel("Importance")
plt.show()

```

Q3. Write a function `obesity_risk_pipeline` to automate the entire pipeline:

Loading and preprocessing the data

Training the model

Evaluating the model

The function should accept the file path and test set size as the input arguments.

```
[37]: # write your function here and then execute this cell
def obesity_risk_pipeline(data_path, test_size=0.2):
    # your code here
    # load data set
    data = pd.read_csv(data_path)
    # standardising continuous numerical features
    continuous_columns= data.select_dtypes(include=["float64"]).columns.tolist()
    scalar = StandardScaler()
    scalar_feature= scalar.fit_transform(data[continuous_columns])

    obesity_risk_pipeline(file_path, test_size=0.2)
```

```
Cell In[37], line 2
def obesity_risk_pipeline(data_path, test_size=0.2):..
```

```
SyntaxError: invalid syntax
```

[Click here for the solution](#)

```
def obesity_risk_pipeline(data_path, test_size=0.2):
```

```

# Load data
data = pd.read_csv(data_path)

# Standardizing continuous numerical features
continuous_columns = data.select_dtypes(include=['float64']).columns.tolist()
scaler = StandardScaler()
scaled_features = scaler.fit_transform(data[continuous_columns])

# Converting to a DataFrame
scaled_df = pd.DataFrame(scaled_features, columns=scaler.get_feature_names_out(continuous_columns))

# Combining with the original dataset
scaled_data = pd.concat([data.drop(columns=continuous_columns), scaled_df], axis=1)

# Identifying categorical columns
categorical_columns = scaled_data.select_dtypes(include=['object']).columns.tolist()
categorical_columns.remove('NObeyesdad') # Exclude target column

# Applying one-hot encoding
encoder = OneHotEncoder(sparse_output=False, drop='first')
encoded_features = encoder.fit_transform(scaled_data[categorical_columns])

# Converting to a DataFrame
encoded_df = pd.DataFrame(encoded_features, columns=encoder.get_feature_names_out(categorical_columns))

# Combining with the original dataset
prepped_data = pd.concat([scaled_data.drop(columns=categorical_columns), encoded_df], axis=1)

# Encoding the target variable
prepped_data['NObeyesdad'] = prepped_data['NObeyesdad'].astype('category').cat.codes

# Preparing final dataset
X = prepped_data.drop('NObeyesdad', axis=1)
y = prepped_data['NObeyesdad']

# Splitting data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size, random_state=random_state)

# Training and evaluation
model = LogisticRegression(multi_class='multinomial', max_iter=1000)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))

# Call the pipeline function with file_path
obesity_risk_pipeline(file_path, test_size=0.2)

```

## 1.6.5 Congratulations! You're ready to move on to your next lesson!

### 1.7 Author

Abishek Gagneja

### Other Contributors

Jeff Grossman

<!-- ## Changelog

Date | Version | Changed by | Change Description |

|:-----|:-----|:-----|:-----|

2024-11-05 | 1.0 Abhishek Gagnejan | Fresh version created |

2025-05-13 | 1.1 Anita Verma | Added the solution code for Ovo model Q2 |

© IBM Corporation. All rights reserved.